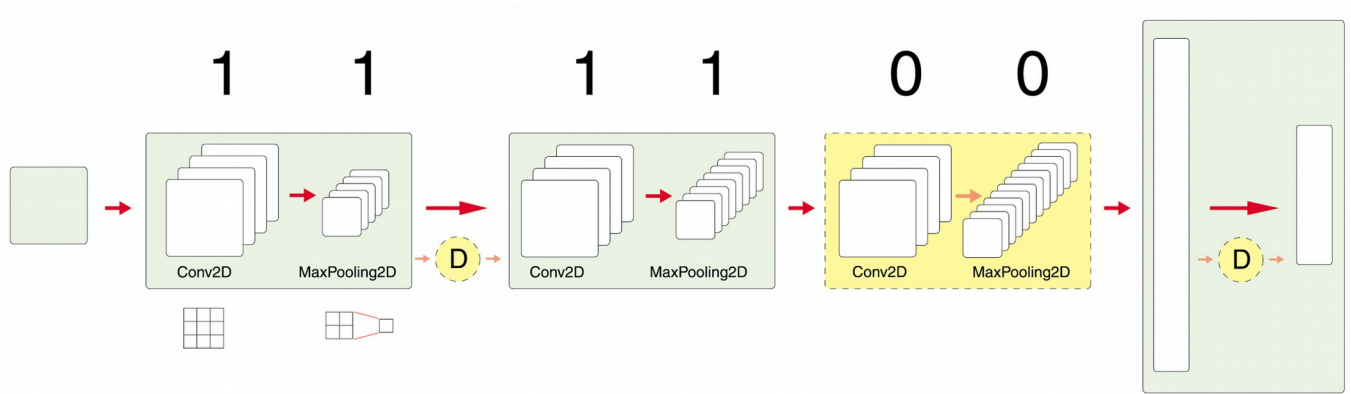


# Report

## Something about Convolutional Neural Networks



In this task, we were provided with a Convolutional Neural Networks previously created using the Tensorflow and Keras library.

Neural Network information is available on request and it is:

```
print(model.summary())

history = model.fit(x_train, y_train, batch_size=6000, epochs = 1, validation_split=0.

model.evaluate(x_test, y_test) #
```

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 26, 26, 32) | 320     |
| max_pooling2d (MaxPooling2D)   | (None, 13, 13, 32) | 0       |
| conv2d_1 (Conv2D)              | (None, 11, 11, 32) | 9248    |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 32)   | 0       |
| flatten (Flatten)              | (None, 800)        | 0       |
| dense (Dense)                  | (None, 128)        | 102528  |
| dense_1 (Dense)                | (None, 10)         | 1290    |
| Total params: 113,386          |                    |         |
| Trainable params: 113,386      |                    |         |
| Non-trainable params: 0        |                    |         |
| None                           |                    |         |

This configuration will be considered basic and we will call the **main model**.

Changing the input parameters to investigate the work of the neural network in terms of computation time, loss function and accuracy.

Datasets to research the model: `tf.keras.datasets.fashion_mnist`

### Experiment 1.

The first experiment was to run the program on various computing platforms.

tab 1

|           | Time(s, ms) | Loss  | Accuracy (%) |
|-----------|-------------|-------|--------------|
| LapTop    | 1583/5      | 0,415 | 84,90%       |
| Colab     | 257/0,9     | 0,42  | 84,40%       |
| Colab+GPU | 33/0,112    | 0,41  | 85,00%       |

### Conclusion 1.

As can be seen from the table, the execution speed in Colab + GPU is fundamentally higher. In the future, all experiments will be carried out on **Colab + GPU**

## Experiment 2.

Differences in the values of loss function and accuracy motivated to conduct one simple experiment. Several runs of the model without changing the input parameters.

tab 2

|                 | Time(s, ms) | Loss  | Accuracy (%) |
|-----------------|-------------|-------|--------------|
| Main model      | 33/0,112    | 0,410 | 86,02%       |
| Main model      | --          | 0,439 | 84,60%       |
| Main model      | --          | 0,428 | 84,10%       |
| Main model      | --          | 0,409 | 85,14%       |
| Main model      | --          | 0,422 | 84,34%       |
| arithmetic mean |             | 0,418 | 84,84%       |

## Conclusion 2

Output parameters are different, but not fundamentally and randomly. For more stable values of all experiments, it seems reasonable to conduct a series of model launches (for example, 10 each) and take into account the arithmetic mean. Using “shuffle=False” in fit-function gives less difference in results also.

## Experiment 3.1

Here we will change the number of feature maps.

Recall in the base model 32 of feature maps.

tab 3.1

| Feature maps | Time(s, ms) | Loss  | Accuracy (%) |
|--------------|-------------|-------|--------------|
| 32 (repeat)  | 33/0,112    | 0,410 | 85,00%       |
| 16           | --          | 0,422 | 84,90%       |
| 64           | --          | 0,421 | 85,00%       |

## Conclusion 3.1

Let's try other values:

## Experiment 3.2

What is the minimum number of features maps with an accuracy of at least 0.8 ?

tab 3.2

| Feature maps | Time(s, ms) | Loss  | Accuracy (%) |
|--------------|-------------|-------|--------------|
| 32 (repeat)  | 33/0,112    | 0,410 | 85,00%       |
| 8            | / 0,09      | 0,450 | 83,40%       |
| 4            | --          | 0,468 | 82,80%       |
| 3            | --          | 0,537 | 79,99%       |
| 2            | --          | 0,558 | 78,00%       |
| 1            | --          | 0,584 | 78,60%       |

**Experiment 4** (series similar experiments):*tab 4.1*

| <b>Feature maps</b>  | <b>Time(s)</b> | <b>Loss</b> | <b>Accuracy (%)</b> |
|--|----------------|-------------|---------------------|
| <b>Main model<br/>1-1-1-1</b>                              | 33             | 0,410       | 85,00%              |
| <b>Remove the final<br/>Conv &amp; Maxpool<br/>1-1-0-0</b> | 32             | 0,388       | 86,18%              |
| <b>Remove the final<br/>Convolution<br/>1-1-0-1</b>        | 31             | 0,418       | 84,97%              |
| <b>Remove the final Maxpool<br/>1-1-1-0</b>                | 35             | 0,362       | 87,15%              |
| <b>Remove<br/>all Conv &amp; Maxpool<br/>0-0-0-0</b>       | 26             | 0,448       | 84,30%              |
| <b>Remove all Convolution<br/>0-1-0-1</b>                  | 25             | 0,726       | 72,71%              |
| <b>Remove all Maxpool<br/>1-0-1-0</b>                      | 36             | 0,329       | 88,14%              |
| <b>Remove the first<br/>Convolution<br/>0-1-1-1</b>        | 31             | 0,488       | 82,45%              |
| <b>Remove the first Maxpool<br/>1-0-1-1</b>                | 37             | 0,328       | 88,35%              |
| <b>adding Convolutions<br/>1-0-1-1 + 1</b>                 | 36             | 0,337       | 87,79%              |
| <b>adding Conv &amp; maxpool<br/>1-0-1-1 + 1-1</b>         | 41             | 0,389       | 86,59%              |

**Add Dropout layer***Tab 4.2*

| <b>Between: 1-0 - D - 1-1</b> | <b>Loss</b> | <b>Accuracy (%)</b> |
|-------------------------------|-------------|---------------------|
| <b>Dropout 0,25</b>           | 0,388       | 86,32%              |
| <b>Dropout 0,5</b>            | 0,411       | 85,66%              |

*Tab 4.3*

| <b>Between full connection:<br/>1-0-1-1_1-D-1</b> | <b>Loss</b> | <b>Accuracy (%)</b> |
|---|-------------|---------------------|
| <b>Dropout 0,25</b>                               | 0,385       | 86,27%              |
| <b>Dropout 0,5</b>                                | 0,382       | 86,02%              |

## Experiment 5:

How many epochs are optimal for maximum accuracy in TEST? (main model 1-1-1-1)

tab 5

| epochs | Time(s, ms) | Loss   | Accuracy (%) |
|--------|-------------|--------|--------------|
| 5      | 33 / 0,1    | 0,3100 | 86,70%       |
| 10     | 61 / 0,1    | 0,2965 | 89,46%       |
| 20     | 120 / 0,1   | 0,2790 | 90,33%       |
| 25     | 150 / 0,1   | 0,2770 | 90,28%       |
| 30     | 180 / 0,1   | 0,2850 | 90,59%       |
| 50     | 300 / 0,1   | 0,3470 | 90,25%       |
| 100    | 610 / 0,1   | 0,4032 | 90,72%       |

## Finally Experiment

feature maps = 32, epochs=25, model 1-0-1-1

tab 6

|                               | Time(s) | Loss test | Accuracy test (%) | Loss train | Acc. train |
|-------------------------------|---------|-----------|-------------------|------------|------------|
| <b>Main model<br/>1-0-1-1</b> | 184     | 0,300     | 90,67%            | 0,13%      | 95,83%     |

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()
x_train = x_train / 255.
x_test = x_test / 255.

x_train = x_train.reshape(60000, 28, 28, 1)
x_test = x_test.reshape(10000, 28, 28, 1)

model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    #tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    #tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    #tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    #tf.keras.layers.Dropout (0.25),
    tf.keras.layers.Dense(10, activation='softmax'),
])

model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
print (model.summary())

early_stopping_callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=2)

history = model.fit(x_train, y_train, epochs = 1, shuffle=False, callbacks=[early_stopping_callback])
model.evaluate(x_test, y_test)
```

| Layer (type)                 | Output Shape       | Param # |
|------------------------------|--------------------|---------|
| conv2d (Conv2D)              | (None, 26, 26, 32) | 320     |
| conv2d_1 (Conv2D)            | (None, 24, 24, 32) | 9248    |
| max_pooling2d (MaxPooling2D) | (None, 12, 12, 32) | 0       |
| flatten (Flatten)            | (None, 4608)       | 0       |
| dense (Dense)                | (None, 128)        | 589952  |
| dense_1 (Dense)              | (None, 10)         | 1290    |
| Total params: 600,810        |                    |         |
| Trainable params: 600,810    |                    |         |
| Non-trainable params: 0      |                    |         |
| None                         |                    |         |

# Implementation callback on the loss function

```
38] (x_train, y_train), (x_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()
x_train = x_train / 255.
x_test = x_test / 255.

x_train = x_train.reshape(60000, 28, 28, 1)
x_test = x_test.reshape(10000, 28, 28, 1)

model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    #tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation='softmax'),
])

model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
#print (model.summary())

early_stopping_callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=2) # 'acc'

history = model.fit(x_train, y_train, epochs = 100, shuffle=False, callbacks=[early_stopping_callback])
model.evaluate(x_test, y_test)
```

---

```
60000/60000 [=====] - 7s 109us/sample - loss: 0.0278 - acc: 0.9910
Epoch 55/100
60000/60000 [=====] - 7s 109us/sample - loss: 0.0249 - acc: 0.9927
Epoch 56/100
60000/60000 [=====] - 7s 109us/sample - loss: 0.0226 - acc: 0.9931
Epoch 57/100
60000/60000 [=====] - 7s 110us/sample - loss: 0.0231 - acc: 0.9929
Epoch 58/100
60000/60000 [=====] - 7s 110us/sample - loss: 0.0246 - acc: 0.9921
10000/10000 [=====] - 1s 89us/sample - loss: 0.6210 - acc: 0.8984
[0.6210185083660995, 0.8984]
```

After the implementation of callbacks and using the splitting of the training set into training and validation set, the optimal number of epochs was 16.

So, one last neural network training was conducted, and the result was obtained:

## Other parameters of neural network:

feature maps = 32, shuffle= True

tab 7

|                               | Epochs    | Time(s) | Loss test    | Acc. test     | Loss train | Acc. train |
|-------------------------------|-----------|---------|--------------|---------------|------------|------------|
| <b>Main model<br/>1-1-1-1</b> | <b>5</b>  | 33      | 0,418        | 84,79%        | 0,410      | 85,00%     |
| <b>Main model<br/>1-0-1-1</b> | <b>20</b> | 184     | <b>0,259</b> | <b>91,10%</b> | 0,155      | 94,36%     |

```
Epoch 20/20
60000/60000 [=====] - 7s 111us/sample - loss: 0.1546 - acc: 0.9436
10000/10000 [=====] - 1s 90us/sample - loss: 0.2587 - acc: 0.9110
```

Out[64]: [0.2587348137915134, 0.911]